

OpenAI Agents SDK Quiz – 48 questions – 2 Hours time.

Q1) Consider the following Python code snippet using the OpenAI Agents SDK:

```
from agents import Agent, function_tool
from agents.extensions.visualization import draw_graph

@function_tool
def process_data(input_str: str) -> str:
    """Processes the input string."""
    return input_str.upper()

agent1 = Agent(name="Processor", tools=[process_data])
agent2 = Agent(name="Finalizer", instructions="Formats the processed data.")
triage_agent = Agent(
    name="Router",
    instructions="Route to Processor then Finalizer.",
    handoffs=[agent1, agent2]
)
# Assume the runner logic is in place to route appropriately
# draw_graph(triage_agent)
```

If `draw_graph(triage_agent)` were called, which of the following accurately describes the visual representation of `agent1` and its connection to `process_data` based on the SDK's agent visualization features?

- a) `agent1` as a green ellipse, connected to `process_data` as a yellow box by a solid arrow.
- b) `agent1` as a yellow box, connected to `process_data` as a green ellipse by a solid arrow.
- c) `agent1` as a green ellipse, connected to `process_data` as a yellow box by a dotted arrow.
- d) `agent1` as a yellow box, connected to `process_data` as a green ellipse by a dotted arrow.

Correct option: D

Q2) In the OpenAI Agents SDK, what is the primary purpose of an Agent's output_type parameter when set to a Pydantic object?

- a) To specify the expected input format for the next agent in a workflow.
- b) To validate the agent's instructions string at runtime.
- c) To enforce that the agent's output is a structured object matching the specified type instead of plain text.
- d) To define the parameters for function tools used by the agent.

Correct option: C

Q3) Consider the following code snippet:

```
from agents import Agent, ModelSettings
agent = Agent(
    name="Configured Agent",
    instructions="Follow instructions precisely.",
    model="gpt-4o",
    model_settings=ModelSettings(temperature=0.5)
)
```

If this agent is used within a Runner.run call where the RunConfig is set with model_settings=ModelSettings(temperature=1.0), what will be the effective temperature used for the LLM calls by this agent?

- a) 0.5, because the agent-specific setting takes precedence.
- b) The temperature is averaged, resulting in 0.75.
- c) 1.0, because settings in RunConfig override agent-specific settings.
- d) The temperature will fluctuate between 0.5 and 1.0 during the run.

Correct option: C

Q4) Which of the following best describes the behavior of the Agent loop when a handoff occurs during Runner.run execution?

- a) The loop terminates and returns the partial result from the current agent.
- b) The original agent attempts to process the handoff target's response.
- c) The current agent and input are updated to the handoff target and its input, and the loop continues.
- d) A new, independent Runner.run process is started for the handoff target.

Correct option: C

Q5) In the context of implementing a custom Function Tool directly (without the `@function_tool` decorator), which of the following is a required component to provide?

- a) A Python function decorated with `@function_tool`.
- b) An async function named `on_invoke_tool` that receives arguments as parsed Python objects.
- c) An async function `on_invoke_tool` that receives the context and arguments as a JSON string, and returns a string output.
- d) A Pydantic BaseModel subclass named `FunctionArgs` to automatically generate the schema.

Correct option: C

Q6) Consider a scenario where you are running a complex workflow involving multiple agents. You want to link all the traces from these agents together as part of a single user conversation thread for easier debugging and monitoring. Which RunConfig parameter is specifically designed for this purpose?

- a) `workflow_name`
- b) `trace_id`
- c) `group_id`
- d) `trace_include_sensitive_data`

Correct option: C

Q7) According to the sources, what is a key constraint regarding the type of context object used within a single agent run?

- a) Each agent, tool function, or lifecycle hook within the run must use a different type of context for isolation.
- b) The context type must be a built-in Python type like dict or str.
- c) The context object must be serializable to be sent to the LLM.
- d) Every agent, tool function, lifecycle, etc., for a given agent run must use the same *type* of context.

Correct option: D

Q8) You are debugging an agent and want to see the raw events streamed directly from the LLM token-by-token to understand exactly how the response is being generated. Which streaming event type should you specifically listen for when using `result.stream_events()` from a `Runner.run_streamed()` call?

- a) AgentUpdatedStreamEvent
- b) RunItemStreamEvent
- c) OutputGuardrailTripwireTriggered
- d) RawResponsesStreamEvent

Correct option: D

Q9) According to the GPT-4.1 Prompting Guide, when providing examples for a tool, where is the recommended placement within the prompt structure?

- a) Within the tool's description field.
- b) In an # Examples section in your system prompt.
- c) Directly within the instructions section of the system prompt.
- d) As comments within the tool's function definition.

Correct option: B

Q10) Consider the Markdown syntax `__*really important*__`. According to the Markdown Guide source, how should this typically be rendered?

- a) This text is **really important**. (Bold only)
- b) This text is *really important*. (Italic only)
- c) This text is ***really important***. (Bold and Italic, using only asterisks)
- d) This text is ***really important***. (Bold and Italic)

Correct option: D

Q11) You have an agent designed to handle customer inquiries. You want to ensure that any question asking for financial advice triggers a specific safety mechanism. Which SDK feature is best suited for this task, applied to the user's input before the agent processes it?

- a) An Output Guardrail
- b) An Input Guardrail with a Tripwire
- c) An Agent as a Tool
- d) A Handoff to a different agent

Correct option: B

Q12) According to the GPT-4.1 Prompting Guide, what is a primary reason to exclusively use the tools field in the OpenAI API request to pass tool definitions, rather than manually injecting tool descriptions into the system prompt?

- a) Manual injection allows for more flexible tool parameter schemas.
- b) Manual injection is the only way to include examples of tool usage.
- c) Using the tools field is the best way to minimize errors and ensure the model remains in distribution during tool-calling trajectories.
- d) Manually injected tool descriptions bypass the need for a separate parser.

Correct option: C

Q13) When configuring the OpenAI Agents SDK, if you do not have an API key from platform.openai.com but still want to run agents using LiteLLM with a different provider, which configuration step is recommended to avoid tracing client errors?

- a) Explicitly set `set_default_openai_key()` to an empty string.
- b) Configure a custom tracing processor that points to `platform.openai.com`.
- c) Disable tracing entirely using `set_tracing_disabled(True)` or set a specific API key for tracing using `set_tracing_export_api_key()` from `platform.openai.com`.
- d) Configure the LiteLLM model to handle tracing internally.

Correct option: C

Q14) Based on the sources, which of the following Markdown syntaxes for creating an ordered list item is NOT recommended for maximum compatibility across different Markdown applications?

- a) 1. First item
- b) 2. Second item
- c) 3. Third item
- d) 1) First item

Correct option: D

Q15) Consider an agent configured with `ModelSettings.tool_choice` set to "required". According to the sources, what happens to this setting automatically after the agent successfully makes a tool call?

- a) It remains "required" to ensure subsequent actions also involve tools.
- b) It is automatically set to "none" to prevent further tool calls in the same turn.
- c) It is automatically reset to "auto" to prevent potential infinite loops.
- d) It is set to the name of the tool that was just called.

Correct option: C

Q16) You are building an agent workflow where a main orchestration agent needs to leverage the capabilities of other specialized agents *without* transferring control to them. Instead, the

orchestrator should treat these specialized agents as functions it can call. Which SDK feature is designed for this specific pattern?

- a) Handoffs
- b) Input Guardrails
- c) Agents as tools
- d) Dynamic instructions

Correct option: C

Q17) Which of the following is NOT a class of tool explicitly mentioned in the OpenAI Agents SDK sources?

- a) Hosted tools
- b) Function calling tools
- c) Agents as tools
- d) Model context protocol (MCP) tools

Correct option: D

Q18) You are implementing a custom input guardrail function. According to the sources, what must this function return?

- a) A boolean indicating whether a tripwire was triggered.
- b) A string containing the reason for the guardrail outcome.
- c) A GuardrailFunctionOutput object.
- d) The modified input that the agent should receive.

Correct option: C

Q19) When using the Agent Visualization feature with Graphviz, what do directed edges represented by solid arrows typically signify?

- a) Tool invocations

- b) Agent-to-agent handoffs
- c) Data flow between components
- d) Errors occurring in the workflow

Correct option: B

Q20) Consider the following code snippet for using dynamic instructions:

```
from agents import Agent, RunContextWrapper

def dynamic_instructions_func(context: RunContextWrapper[UserInfo], agent:
Agent[UserInfo]) -> str:
    # UserInfo is a custom dataclass for context
    return f"Help the user named {context.context.name}."

agent = Agent[UserInfo](
    name="Personalized Assistant",
    instructions=dynamic_instructions_func
)
```

What is the key benefit demonstrated by using a function for the instructions parameter compared to providing a static string?

- a) It allows the agent to use specific tools defined within the function.
- b) It makes the agent's instructions immutable after creation.
- c) It enables the agent's behavior to be tailored based on the current run context.
- d) It automatically handles handoffs to other agents.

Correct option: C

Q21) Which of the following statements about the context object in the OpenAI Agents SDK is accurate based on the sources?

- a) It is automatically populated by the SDK with the full conversation history.

- b) It is primarily used to pass data to the LLM during API calls.
- c) It is a local object that can be used for contextual data, dependencies, and helper functions within the run.
- d) Its type can change dynamically during a single agent run based on the current agent or tool.

Correct option: C

Q22) According to the sources, when building agentic workflows with GPT-4.1, including a specific reminder in the system prompt is highly recommended to prevent the model from prematurely ending its turn before the user's query is fully resolved. What is this type of reminder called?

- a) Tool-calling
- b) Planning
- c) Persistence
- d) Instruction Following

Correct option: C

Q23) You are implementing an agent that requires the LLM to *always* attempt to use a tool if available, but allows the model to choose which tool is most appropriate. How should you configure the agent's ModelSettings.tool_choice parameter to achieve this?

- a) "auto"
- b) "required"
- c) "none"
- d) The specific name of the tool to use.

Correct option: B

Q24) Consider the following simplified agent loop logic described in the sources:

1. Call LLM with current agent/input.
2. LLM produces output.

3. If output is X, loop ends, return result.
4. If output is Y, update agent/input, re-run loop.
5. If output is Z, run tool calls, append results, re-run loop.

Which LLM output corresponds to Y in this description?

- a) final_output
- b) Tool calls
- c) Handoff
- d) Structured output object

Correct option: C

Q25) You are running an agent workflow and receive a `ModelBehaviorError`. According to the sources, what does this exception typically indicate?

- a) An error in the code you wrote using the SDK.
- b) The maximum number of turns for the run was exceeded.
- c) A guardrail's tripwire was triggered.
- d) The model produced invalid outputs, such as malformed JSON or attempting to use non-existent tools.

Correct option: D

Q26) When using `@function_tool` on a Python function, how does the Agents SDK typically determine the JSON schema for the tool's parameters?

- a) It uses type hints from the function signature and infers the schema from the function's return type.
- b) It requires the schema to be manually provided via a `params_json_schema` argument in the `@function_tool` decorator.
- c) It automatically creates the schema from the function's arguments, leveraging Python's `inspect`, `griffe` (for docstrings), and `pydantic`.
- d) It prompts the LLM to generate the schema based on the function's name and docstring.

Correct option: C

Q27) According to the GPT-4.1 Prompting Guide's advice on delimiters for long context, which format performed particularly poorly in their testing?

- a) XML
- b) Lee et al.'s format (e.g., ID: 1 | TITLE: ... | CONTENT: ...)
- c) JSON
- d) Pseudo-XML

Correct option: C

Q28) Consider an agent configured with a list of potential handoff agents. The agent's instructions guide it to delegate to the appropriate sub-agent based on the user's request. Which of the following SDK concepts represents the directed edges from the main agent to its potential handoff targets in the agent visualization?

- a) Dotted arrows
- b) Green ellipses
- c) Yellow boxes
- d) Solid arrows

Correct option: D

Q29) You are developing an application using the OpenAI Agents SDK and want to enable debug logging for both the core SDK and the tracing module. According to the sources, what are the names of the Python loggers you should configure?

- a) "agents" and "tracing"
- b) "openai.agents.core" and "openai.agents.tracing"
- c) "agents.core" and "agents.tracing"
- d) "openai.agents" and "openai.agents.tracing"

Correct option: D

Q30) According to the GPT-4.1 Prompting Guide, if conflicting instructions are present in a system prompt, which instruction does GPT-4.1 tend to follow?

- a) The instruction that is more specific.
- b) The instruction that uses stronger emphasis (e.g., all-caps).
- c) The instruction that is closer to the end of the prompt.
- d) The instruction that is listed first in the prompt.

Correct option: C

Q31) When using `Runner.run_streamed()`, the `result.stream_events()` provides different types of `StreamEvent` objects. Which event type is most useful for receiving updates at the level of a completed "message generated" or "tool ran"?

- a) `RawResponsesStreamEvent`
- b) `AgentUpdatedStreamEvent`
- c) `RunContextStreamEvent`
- d) `RunItemStreamEvent`

Correct option: D

Q32) Consider a scenario where you want to run an agent synchronously and retrieve its final output as a `RunResult`. Which method from the `Runner` class should you use?

- a) `run()`
- b) `run_sync()`
- c) `run_streamed()`
- d) `run_async()`

Correct option: B

Q33) According to the sources, what is a potential drawback of prompting GPT-4.1 agents to plan explicitly between tool calls using techniques like Chain-of-Thought?

- a) It can decrease the likelihood of the agent using its tools effectively.
- b) It often leads to the agent hallucinating tool inputs.
- c) It can result in higher cost and latency due to using more output tokens.
- d) It makes the agent less steerable and responsive to instructions.

Correct option: C

Q34) When using `@function_tool` to wrap a Python function, what is the expected input type for the function's arguments by default, assuming standard type hints are used?

- a) They must be JSON strings.
- b) They must be Python primitive types (str, int, float, bool).
- c) They must be Pydantic BaseModel instances.
- d) They can be any Python type that can be wrapped in a Pydantic TypeAdapter, with schema automatically generated.

Correct option: D

Q35) According to the Markdown Guide, how should you create an unordered list using Markdown syntax?

- a) Using numbers followed by periods (e.g., 1. Item)
- b) Using letters followed by periods (e.g., a. Item)
- c) Using any of -, *, or + followed by a space.
- d) Using angle brackets (e.g., <item>)

Correct option: C

Q36) You are configuring an agent and want to prevent it from using any of the tools provided to it, regardless of the user input. How should you set the `ModelSettings.tool_choice` parameter?

- a) "auto"
- b) "required"
- c) "none"
- d) An empty list [].

Correct option: C

Q37) Which of the following correctly describes the behavior when a `MaxTurnsExceeded` exception is raised during `Runner.run`?

- a) It indicates the agent successfully reached a final output after many turns.
- b) It is raised when the LLM produces invalid output.
- c) It is raised when a guardrail tripwire is triggered.
- d) It is raised when the agent loop runs for more turns than allowed by the `max_turns` parameter.

Correct option: D

Q38) According to the GPT-4.1 Prompting Guide, what is a potential common failure mode if you instruct a model to *always* follow a specific behavior, such as "you must call a tool before responding to the user," without additional nuance?

- a) The model might refuse to respond at all.
- b) The model may hallucinate tool inputs or call the tool with null values if it lacks sufficient information.
- c) The model might ignore all other instructions in the prompt.
- d) The model will enter an infinite loop of tool calls.

Correct option: B

Q39) In the V4A diff format used by the `apply_patch` tool described in the sources, which symbol is used to precede lines of code that should be added to the file?

- a) –

- b) +
- c) (space)
- d) @@

Correct option: B

Q40) According to the Markdown Guide, what is the recommended method for italicizing text in the middle of a word for emphasis, ensuring compatibility across different Markdown applications?

- a) A_cat_meow
- b) A/cat/meow
- c) A~cat~meow
- d) A*cat*meow

Correct option: D

Q41) When orchestrating multiple agents in the SDK using Handoffs, what is a recommended practice to help LLMs understand the concept of delegation?

- a) Define complex input_filters for every handoff.
- b) Ensure each handoff agent has a unique output_type.
- c) Include information about handoffs in your agents' instructions, potentially using recommended prompt prefixes.
- d) Avoid defining handoffs on the starting agent and only define them dynamically.

Correct option: C

Q42) The RunConfig object allows for several global configurations for a run. Which of the following configurations applies a filter to the inputs passed during *all* handoffs within that run, unless a specific handoff already has its own filter?

- a) input_guardrails
- b) output_guardrails

- c) handoff_input_filter
- d) trace_metadata

Correct option: C

Q43) Which of the following is a limitation of using the convenience method `agent.as_tool()` compared to implementing an agent-as-tool pattern using a custom function tool that calls `Runner.run` directly?

- a) `agent.as_tool()` does not allow specifying tool descriptions.
- b) `agent.as_tool()` does not support setting all configuration options, such as `max_turns` for the tool-agent's run.
- c) `agent.as_tool()` can only be used with hosted tools.
- d) `agent.as_tool()` prevents the tool-agent from using its own set of tools.

Correct option: B

Q44) According to the GPT-4.1 Prompting Guide, what is a recommended strategy if you observe that an agent, built with GPT-4.1, is being too verbose or adding unnecessary formatting in its responses?

- a) Add an # Examples section showing concise outputs.
- b) Use all-caps in the system prompt to strongly discourage verbosity.
- c) Implement an input guardrail to filter out verbose sections of the user's query.
- d) Provide explicit instructions and potentially examples demonstrating the desired response format and tone.

Correct option: D

Q45) When designing function tools with `@function_tool`, the sources mention that you can optionally include the run context as the first argument. What is the type hint you should use for this context parameter?

- a) `RunContext`

- b) RunContextWrapper
- c) RunContextWrapper[YourContextType]
- d) Dict[str, Any]

Correct option: C

Q46) Consider a scenario where you are using a non-OpenAI model provider via LiteLLM that does not fully support the OpenAI Responses API, but does support the Chat Completions API. What is a valid approach mentioned in the sources to resolve potential 404 or similar errors related to API support?

- a) Modify the LiteLLM configuration to emulate Responses API support.
- b) Call `set_default_openai_api("chat_completions")`.
- c) Ensure the non-OpenAI provider explicitly lists Responses API compatibility.
- d) Remove all tools and structured output types from your agents, as these require the Responses API.

Correct option: B

Q47) When using the Agent Visualization feature, which element is represented as an ellipse with a green fill?

- a) An Agent
- b) A Handoff
- c) A Tool
- d) The start node

Correct option: C

Q48) According to the GPT-4.1 Prompting Guide, when using the `apply_patch` tool with the V4A diff format, how are lines of code provided for context around a change snippet?

- a) Using line numbers relative to the file beginning.
- b) Using line numbers relative to the start of the code block being patched.

c) By including a specified number of lines (default 3) immediately above and below the changed code, potentially with @@ markers for class/function context.

d) Context is automatically inferred by the tool and not included in the patch text.

Correct option: C
